Syed Monis Sarwar

Lucknow, India | +91 93367 54624 | monissms16@gmail.com
m0nis.com | github.com/MonisMS | linkedin.com/in/syed-monis-sarwar-sms47

Education

Babu Banarasi Das Institute of Technology and Management — Lucknow, India
B.Tech in Computer Science and Engineering — *2024 – 2028*

Anand Bhawan School — Barabanki, India
ISC (Class XII), 78% — *2023*
ICSE (Class X), 93% — *2021*

Open Source

corsairdev/corsair | *YC W25, 13k+ stars, TypeScript* — PR #703, merged
- Added the OCR.space integration as a standalone typed package (client, endpoints, error handling, test suite), reviewed and merged by the founder

bholmesdev/hubble.md | *1.4k+ stars, Markdown editor* — PR #241, merged
- Fixed the editor's Enter handler dropping users out of a list when pressing Enter after an image in a list item

twitbruv/twitbruv | *TypeScript* — PR #113, merged
- Handled @-prefixed queries in user search and added a fallback for avatars that fail to load

Projects

APIx: Daily Airfare Price Index for India | *PostgreSQL, Next.js, GitHub Actions* — Live | GitHub
- Built a daily collection pipeline on GitHub Actions that prices 12 DGCA-weighted routes across 5 advance-purchase windows, storing **24,448 fares** in Postgres
- Implemented the official CPI method (Jevons elementary indices, Young aggregation, outlier cleaning, imputation), **reproducing MoSPI's published worked examples to 4 decimal places**
- Made every number auditable: each index point carries a hash of its inputs, and each publish writes a full new version in one transaction so readers never see half-written data
- Served the index through a public JSON API that falls back to a static snapshot when the database is down

DataPilot: Prompt-to-Dataset LLM Agent | *TypeScript, Next.js, PostgreSQL, LLM APIs* — Live | GitHub
- Built an agent pipeline where an LLM planner turns a plain-English request into a schema and plan, runs **9 source connectors in parallel**, and extracts rows in batched LLM calls
- Guarded against hallucinated rows by requiring a source quote per row and checking it against the fetched page in code, then validating types, merging duplicates and scoring each row
- Added an SSRF guard for fetched URLs (private, loopback and cloud-metadata IPs) and fenced scraped text as data to resist prompt injection
- Wrote a provider-agnostic LLM client (Gemini, Anthropic, OpenAI) with model fallbacks and a rule-based mode, plus an eval harness scoring fill rate and precision

Curio: Personalized Reading Feed | *TypeScript, Next.js, PostgreSQL, Drizzle ORM, Resend* — Live | GitHub
- Ingested **275+ sources across 26 topics** (RSS, Hacker News, Reddit, YouTube) in parallel with an 8-second per-source timeout, so one slow feed never stalls a run
- Ranked feeds by blending recency with a per-source quality score from bookmark and read rates; built search on Postgres full-text search instead of an external service
- Automated 5 scheduled jobs (ingestion, digests, quality scoring, cleanup, re-engagement) and sent timezone-aware digests with HMAC-signed unsubscribe links

Technical Skills

Languages: TypeScript, JavaScript, SQL
Frameworks: Next.js, React, Node.js, Express, Tailwind CSS, Zod
Databases: PostgreSQL, Drizzle ORM
AI: LLM APIs (OpenAI, Anthropic, Gemini), tool calling, structured output, LLM evals
Tools: Git, GitHub Actions, Docker, Vercel

Achievements & Leadership

1st place, Smart India Hackathon 2026 internal round, BBDITM: APIx (problem SIH26056, MoSPI)
1st place, BBDITM Internal Hackathon 2026: Inner Council, a voice-first guidance agent
Campus Lead, Google Developer Groups (GDG) on Campus, BBDITM — Oct 2026 – Present
Campus Lead, AlgoZenith, BBDITM — Aug 2025 – Aug 2026